

Ph.D. Statement of Purpose

section: introduction

TODO: a good summary of the applicant's background, interests, and motivation for doing a PhD

- who are you as a researcher?
- what drives You?
- what do you want to research and accomplish?

Despite the impressive capabilities of current AI models, leveraging AI in human-AI teams and human-centered task is yet an open problem, one I am deeply interested in. As a predoctoral researcher at University of Chicago, my research focuses on **human-centered machine learning (HCML)**; in particular, we are designing AI-assistance frameworks for medical training. My research and experience had inspired my interests in the following problems: (1) designing AI assistance that inspires humans' appropriate reliance instead of blind trust (2) AI learning human perception to provide better AI assistance .

1. How can we design AI assistance that inspires humans' appropriate reliance?

Many human-AI teams provide humans with "persuasive" AI assistance that improves humans performance but at the expense of humans' overreliance and blind trust towards AI [1]. This is undesired and unethical, especially in high-stake domains where humans should have the last say. On the other hand, many work including ours also show that neutral, non-persuasive assistance that dissuade humans from blindly following AI perform no better than persuasive assistance. Thus, I am interested in solving this dilemma.

section: research exp

I worked with Prof. Chenhao Tan at UChicago to design a human-compatible model for case-based decision support.[2]. Our model learns both classification and predicting human perception; we conducted synthetic experiments and human studies on natural image and medical dataset, showing our human-compatible representation leads to better decision support performance than traditional AI representations.

section: research experience

subsection: research experience

At the Pre-Doctoral Masters Program in Computer Science at University of Chicago, I worked with Dr. Chenhao Tan on designing AI-driven decision support and training systems for medical diagnosis. Motivated by AI learning human intuition, we devised a human-compatible model that learned both a classification task and predicting human perception. Such a human-compatible representation learned some form of human similarity function and proved to be useful for case-based decision support.

subsection: challenges and contributions

Our work was a unique project with no preexisting framework to build on, so we face many unprecedented problems. On the data side, we had to deal with issues on the format of human perception data, inter-annotator agreement, the amount of data to collect, etc; on the model side, we had to decide on our model architecture, which layer's embedding to use, the dimension of the embedding, etc.

As human studies were expensive, what helped in forming our decisions were synthetic experiments. I built a synthetic dataset of digitally-generated insects with controllable features and also simulated human agents by tuning weights on the features; I also generated varying decision boundaries by altering the ground-truth labels. Synthetic experiments allowed us to exhaustively explore our design and hyperparameters and also provided more insight; for example, our method works better for non-linearly separable distributions.

With positive results from synthetic experiments, we moved on human studies. I conducted a human study on a chest X-ray dataset with Prolific crowdworkers. Compared to an AI that only learned classification, our human-compatible representations provided decision support that led to better performance pneumonia diagnosis.

subsection: new skills

Of the many research skills I learned, the most important skill to me was using experiments to verify ideas. This was counter-intuitive to me at first since I was used to learning the in-and-outs of a well-rounded idea first and then getting hands-on with it. But our project was a nascent one where we had to build everything from scratch; whenever we had the smallest idea we would run experiments, often using simulated human agents. This eventually led to our "hunch" that a model learning human perception could be useful in providing decision support.

Another skill is to be flexible and be willing to reshape the problem to be solved. Our initial goal was to design an AI-driven tutorial for radiology residents. In thinking about how to model human learners, we came across literature on AI learning human perception. With many small, incremental ideas and experiments, we found that AI learning human perception produced representations more aligned with humans. But we did not know how this could be useful for teaching, so we detoured to an easier problem of case-based decision support: providing assistance as the test case's nearest-neighbor in the training set using our learned similarity function. After showing our method works for case-based decision support, we return to the problem of teaching humans, with much more foundation to build on.

subsection: future directions

Our experiment results showed our human-compatible representation leads to better decision support performance than an AI that only learns classification. While the results are positive, many issues remain unsolved in this nascent direction. This research project inspired my interest in the following problems:

(i) generating AI assistance that inspires appropriate trust and reliance on the AI. Many past work provides AI's decision and explanation as assistance and show improved human-AI team performance, but a great part of the improved performance can be attributed to humans simply following AI's suggestion, suggesting overreliance on the AI and a lack of human agency. (citations) Our decision support framework instead aims to provide neutral support that retains as much human agency as possible: we provide example explanations from each class and do not reveal AI's predicted label. However, our experiments also shows providing neutral support leads to lower performance than support that nudges human to follow AI's decision. That is, in human-AI teams there is a tradeoff

between human agency and task performance. I am interested in solving this problem as I believe human agency is an important but underexplored issue in human-AI collaboration.

TODO: potential solutions/proposals to the problem (ii) using AI to learn and model human perception and intuition. As our work showed, AI learning human perception can provide better decisions support performance; I believe in general, incorporating more human knowledge and intuition to AI frameworks has potential to provide more effective AI assistance, but many issues remain unsolved in this nascent direction. For example, our work focuses on visual perception on images as it is the easiest form of perception for humans; perception for other modalities like audio and text are also important but much more difficult to collect and utilize. The limitations and bounds of AI learning human perception is also unexplored, including the amount of data needed, the format of data (we used triplet annotations), how well neural networks can learn, etc.

TODO: potential solutions/proposals to the problem

section: my fit with the school/program

TODO:

- whichs profs?
- projects or research directions?

I want to pursue a Ph.D. in human-centered AI because of I enjoy solving problems in human-AI interaction and aspire to do research that has direct social impact. I believe the research at CMU HCII aligns with my interest and motivation.

TODO:

- sherry wu
- ken
- haiyi zhu

References

- [1] Gagan Bansal, Tongshuang Wu, Joyce Zhou, Raymond Fok, Besmira Nushi, Ece Kamar, Marco Tulio Ribeiro, and Daniel Weld. Does the whole exceed its parts? the effect of ai explanations on complementary team performance. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, CHI '21, New York, NY, USA, 2021. Association for Computing Machinery.
- [2] Han Liu, **Yizhou Tian**, Chacha Chen, Shi Feng, Yuxin Chen, and Chenhao Tan. Towards effective case-based decision support with human-compatible representations. In *Proceedings of the 1st ICML 2022 Workshop on Human-Machine Collaboration and Teaming (HMCaT, ICML 2022)*, 2022.